

Final Year Project: A Computing System in VHDL.

Student: Richard James Howe

May 3, 2013

Abstract

The goal of this project is to create a computing system in VHDL from the ground up in order to make a product that is useful for both teaching and eventually much more. This project includes the firmware and the toolchain that is to target the device.

Keywords: *VHDL, FORTH, SoC, CPU, Assembly.*

Supervisor: Marc Eberhard.

Degree: BSc, Electronic Engineering.

Department: School of Engineering and Applied Science.

Institution: Aston University.

Contents

1 Synopsis	4
1.1 The Project	4
1.2 Structure of this report	4
2 Introduction	4
2.1 Contact details and licenses	5
3 Project Goals	5
3.1 What makes this an educational computer?	5
4 Tools used	5
4.1 Tools list	5
5 The Hardware	6
6 VHDL	7
6.1 J1	7
6.2 Basic system overview	7
6.3 H2	7
6.3.1 Why a stack machine?	7
6.3.2 The H2 CPU	7
6.4 VGA	10
6.5 UART	11
6.6 RAM and other inferred modules	11
6.7 Top level	11
6.8 Test benches and waveforms	12
6.9 Resource usage	14
7 Problems	15
7.1 Build time	15
7.2 Debugging	16
8 Documentation	16
8.1 FORTH interpreter	16
8.1.1 The C Program	17
8.1.2 Basic commands and ideas	20
8.2 H2 CPU	21
8.2.1 Assembly	21
9 Future plans	23
9.1 Hardware	23
9.1.1 Different platforms	23
9.1.2 More generic code	24
9.1.3 PS/2 or USB keyboard	24
9.1.4 Signal Generator	24
9.1.5 Data logger	24
9.1.6 Logic Analyser	25
9.1.7 Multiple cores	25
9.1.8 Miscellaneous	25

9.2	Firmware	25
9.3	Build system	25
10	Conclusion	26
10.1	Project achievements	26
10.2	What I would have done differently	26
10.3	Final words	27
11	References	27
12	Appendix	28
12.1	H2.VHD	28
12.2	Original specification	33
12.2.1	FPGA: Main plan.	33

List of Figures

1	ALU Operation(T Next)	9
2	ALU Other	9
3	Instruction set encoding.	9
4	The H2 CPU Core	10
5	VGA test pattern	11
6	Adder waveform	12
7	Adder waveform	13

1 Synopsis

Here is a brief synopsis on the project, while containing redundant information will give a brief overview of the project for those not interested in reading the entire report.

1.1 The Project

The idea of this project is to create an educational computer for electrical engineering students being taught VHDL and to eventually extend this project so that they could use this project as a piece of lab equipment which is outside of the scope of this project.

The goals of this project are as follows:

- Develop a tool chain for a CPU core.
- Create a CPU in VHDL that is designed to execute FORTH code efficiently.
- Make peripherals to go around this device and interface them with it.
- Develop simple firmware to reside on this device.

All of those goals have been largely met and I will continue to work on this project after university.

The project also has several *design* goals:

- **Code Portability:** *All* code should be written as to be portable.
- **Resource Efficient:** The project should not use up a lot of resources on the device.
- **Extensibility:** The code should be extensible and modular.

This project will target a development board called the **Nexys 3**, an FPGA development board from Digilent.

1.2 Structure of this report

The report will detail the workings of the project, document how to run it and recreate the system, talk about design decisions and then outline my future plans, that is where I would like to take this system in the future.

Documentation is going to be a big feature of this

2 Introduction

The idea of this project is to create an educational system for electrical engineers who are studying VHDL that can be eventually developed into something more useful, and still have those engineers in mind.

This is meant to be an entire system, which is why the project spans multiple languages and includes different sub-projects, which can be put into roughly three different fields: The C/FORTH assembler program (which also finds other

uses), the Assembler that is to run on the device and finally the VHDL that implements the device itself.

I have used other peoples modules in this project and that will be clearly labelled, the intention was to use them and then swap them out, but some have increased functionality above the original goals and there were time constraints as well.

2.1 Contact details and licenses

All code, including this thesis, has been released under an open-source license (LGPL).

The project is available on Github here: <https://github.com/howerj/fyp.git>. Please note that as I am working on this project continually to see the version that this thesis relates to please check out tag *v0.7*.

I intend to continue working on the project after university improving functionality, rewriting sections and porting to different devices.

3 Project Goals

The goal of this project is to create an educational computer, it is as much as for the education of other people as it was for me, I do not expect it to be used as a prototype for a real system until a few years of continual development.

3.1 What makes this an educational computer?

Nothing by itself makes this system educational, a course would have to be arranged around the device. Perhaps lessons could be given to create a new module for the system, as in the section 9 (Future plans). Alternatively they could be given the task to write software for the computer, or even a compiler for the system which is more of a computer science project. A C compiler for a subset of C would be potentially very useful and possibly within the scope of a final year computer science student.

A module could be removed from the project which the student would then have to implement for a more gentler introduction to VHDL.

But the real use would be when the additional modules module are complete, again described in 9. This would provide an engineering student with a working piece of lab equipment that they could either use in class, or for the more ambitious at home.

4 Tools used

As this project is entirely software based you will need a list of all the tools I have chosen, that will be included in this section as well as why I have used these tools.

4.1 Tools list

As of the 27th of April, 2013, I have used the following to run and develop my project:

- Debian 6.0 (This includes a lot of the software used, eg. Gcc)
- Xilinx Webpack ISE 14.2 (Free for students).
- Git, A distributed version control system.
- GHDL, Digital simulation for VHDL.
- GTKWave, Waveform viewer.
- Make, For the VHDL build process.
- Gcc, The GNU C compiler.
- Bash, command interpreter.
- Digilent's programmer for the Nexys 3 device.

5 The Hardware

The only pieces of hardware need are; a laptop, a VGA capable monitor, two micro USB cables, a VGA cable and Nexys 3 development board available from Digilent [1].

The Nexys 3 board forms the core of this project, it is the device that I will be targeting. Seeing as this is an external block that is provided as is there is no real need to go into too many details. It provides an FPGA that has plenty of room for my project (in terms of Look Up Tables, Configurable Logic Blocks, Block RAMs, etcetera) as well as nice interfaces for external hardware (USB UART, The device can be programmed over USB, a VGA port, LEDs, Switches).

The FPGA I will be using is called the XC6LX16-CS324, part of the Spartan-6 family from Xilinx[2]. While the datasheet provides all the relevant details a quick overview would not go amiss. The device provides for; 12 * 18kB Block RAM (dual port), 32 Digital Signal Processing "slices"¹ (which includes an adder and a multiplier) and roughly 2300 Configurable Logic Blocks (CLBs). CLBs are the bread and butter of the FPGAs, they are not a canonical device² such as an adder or an 'OR' gate, but each FPGA is comprised of a version of this.

An FPGA consists of an array of these devices; CLBS and miscellaneous other pieces of hardware such as the DSPs and BRAMs mentioned, each CLB has a Look Up Table (or multiple ones) which must be configured, this describes how the logic behaves and some flip flops for holding the internal state. These blocks must then be routed into something useful, the places and routing of these resources is a difficult problem computationally and takes up a long time.

These internals are connected to the outside world by a set of IOBs (Input/Output Blocks) which provide a configurable interface for each physical pin on the IC package.

More information on CLBs and IOBs for Xilinx devices can be found in the references[3].

¹A term used by Xilinx to denote a physical on chip device

²Some vendors call CLBs by a different name but they have *roughly* the same function

6 VHDL

The main thrust of this project lies in the VHDL, all of this project revolves around the architecture defined here. Although most of the assembler was created separately as a fully blown language, the definitions for the instructions are dependant on what is going on in the CPU core naturally.

6.1 J1

The project is built around a translation and improvement of the J1 core [4], a small stack processor built in Verilog and optimized to efficiently execute FORTH instructions, most of which can be executed in one clock cycle. It was perfectly suited for my project although it was not written in my language of choice.

6.2 Basic system overview

A basic overview of the system is as follows. First we write an assembly program

6.3 H2

The original core was written in Verilog, a language that lends itself to more compact code and where things are not as explicit, for example the exact size of variables.

Getting the translation to work was fairly challenging and I also experimented with the core more, I moved some instructions around to allow for more ALU operations and added a few more instructions.

I renamed this core the "**H2**" and it will be referred to by either this name or as "**CPU**" unless I say it refers to another one.

6.3.1 Why a stack machine?

Stack machines have several advantages when it comes to embedded development compared to the more mainstream register machines, naturally they have disadvantages which also will be addressed, but they do not really apply here.

This architecture tends to produce denser code[5], this is because the operands tend to be implicit due to the fact that most operations happen on the stack. You do not need to specify where each instruction is to get it's operands as in a register based computer. This code is further reduced in size (potentially) by the execution model espoused by FORTH systems which execute threaded code [6].

Other advantages stack machines have include minimal processor state, the H2 only has a program counter and two stacks and fast interrupt response time³.

6.3.2 The H2 CPU

The stack based H2 CPU is the core of this project around which all else is built, it is a stack machine as said, it has a fairly simple architecture which I will describe.

³Interrupts have not been implemented as of yet

As this core is so central to the project I am including it in the appendix for the records.

The H2 has two stacks, a return and a data stack which are 32 and 33 machine words deep respectively. It operates on 16-bit values, but with minor modifications could be made to work on greater (but not smaller) bit widths. The instruction set encoding is fairly dense and uses the full sixteen bits although there is still some room for extra ALU instructions. To speed the system up dual port RAM is used, one instruction can be issued and completed in one clock cycle. A Von Neumann architecture is used, however that is not actually dependant on the CPU itself but how it is wired up to the RAM.

A list of instructions is given bellow, on a separate page to act as a handy reference:

Number	Instruction
0	T
1	N
2	$T + N$
3	T and N
4	T or N
5	$T \oplus N$
6	not T
7	$T = N$
8	$N < T$ (signed)
9	N logical right shift T
10	$T - 1$
11	R
12	[T]
13	N logical left shift T
14	depth
15	$N < T$ (unsigned)
16	$T - N$
17	not ($T \oplus N$)
18	(reserved for multiplication)
19	(reserved for multiplication)
20	$T \leftarrow \text{Input}$
21	Write output
22	N rotated right by T
23	N rotated left by T
24	Clear T
25	Reserved
...	...
31	Reserved

Figure 1: ALU Operation(**T Next**)

Field	Width	Action
T Next	5	ALU Operation (may replace T)
$T \rightarrow N$	1	Copy T to N.
$T \rightarrow R$	1	Copy T to R.
$N \rightarrow [T]$	1	RAM write.
$R \rightarrow \text{PC}$	1	Copy R to PC
dstack \pm	2	Signed increment of data stack.
rstack \pm	2	Signed increment of return stack.

Figure 2: ALU Other

The H2 Instruction set is very densely packed, it is described in these three tables. ALU operations can happen in conjunction with RAM and input/output read/writes.

Instruction	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Literal	1	Literal Value														
Jump	0	0	0	Address												
Conditional Jump	0	0	1	Address												
Call	0	1	0	Address												
ALU	0	1	1	T Next			$N \uparrow$		$T \uparrow$	$R \uparrow$	$[T] \uparrow$	$R \uparrow \text{PC}$	rstack \pm		dstack \pm	

Figure 3: Instruction set encoding.

The above figures completely describe the instruction set for the H2 CPU, every single bit is occupied and if an ALU instruction happens then there are many multiple options as to what can be done, read and writes can be done simultaneously with arithmetic and stack operations speeding certain tasks up.

A diagram of the flows of data is shown below, the BRAM on the left can be seen as the current stack and the one on the right as the next state:

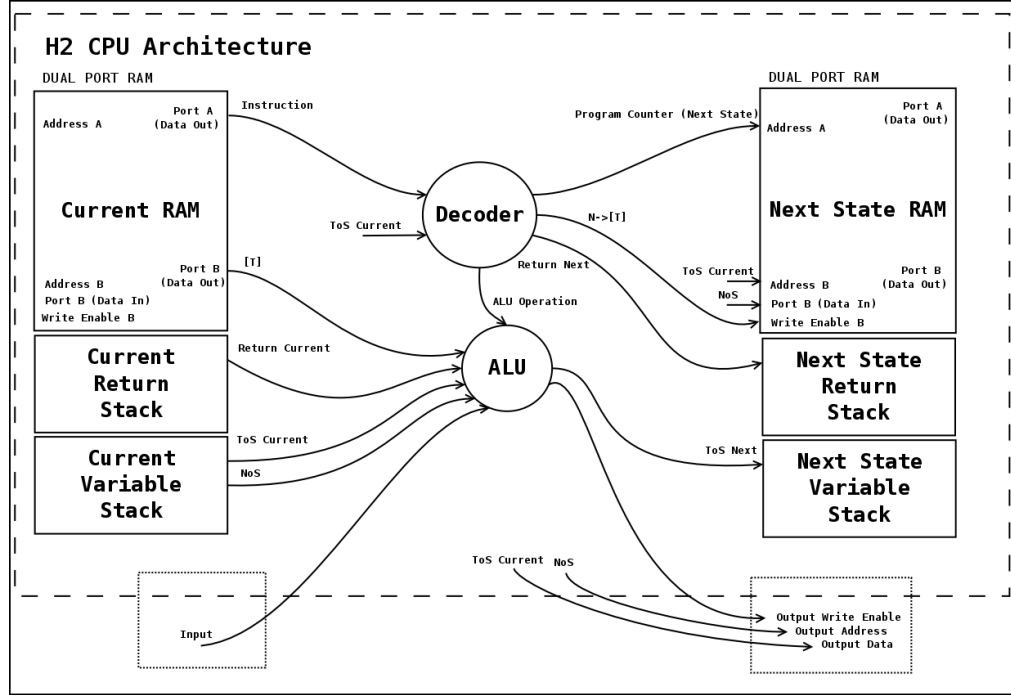


Figure 4: The H2 CPU Core

Most FORTH instructions can be executed in a single clock cycle, for example to add two numbers three separate instructions have to be run simultaneously for it to happen; an ALU $T+N$, $N \rightarrow T$ and a data stack delta of -1 .

Jump indirection can be performed, however it would take multiple clock cycles to do so; First any calculations and fetching of data for the address have to be done which would take an indeterminate amount of time depending on what you want to do, then the data would have to be moved to the return stack ($T \rightarrow R$) and then after this $R \rightarrow PC$ along with the relevant stack deltas. This takes a minimum of three cycles making it more expensive to do.

6.4 VGA

The VGA module acquired online at [7] provides an 80 column by 40 row text buffer with a resolution of 640 by 480. It is simple to interface with and required only minor modification to get up and running.

A picture of the output can be seen here with a test pattern loaded into the text buffer RAM:



Figure 5: VGA test pattern

To interface with this two BRAMs are initialized, one with the test pattern, the other with a font (the font BRAM actually acts as a ROM, but internally they are the same device). To use the unit there are only a few controls, a single control register (on/off, colour output, etcetera), two cursor positioning registers and the BRAM text buffer, simply write what you want to that buffer and it will be displayed on the screen.

6.5 UART

The UART is the primary method of talking to the device, it is customizable at compile time where the baud rate can be set. This module was acquired online and is in the reference[8]. It is a standard UART core that is interfaced to the H2 in the top level. It is capable of both transmission and reception. I will want to rewrite this section in the future to bring it into line with exactly what I would like, it currently gives many warnings when synthesized, none of which cause any problems at the moment.

6.6 RAM and other inferred modules

Instead of using Xilinx specific components you can use VHDL to create code that fits a certain template, this template is picked up by the synthesizer and you can then verify if this is what you intended. Using this instead of instantiating proprietary blocks and loading the RAM with proprietary tools is a much better way of doing things that allows me more flexibility in how I go about things.

The VGA unit and the CPU both have their own RAM blocks, two for the former and one for the latter. They are initialized from an ASCII encoded binary format, which while not efficient is certainly easy to process with the given tools.

For a given file describing the RAM, for example "**h2_mem.vhd**", the initial contents will be described in a file called "**h2_mem.binary**".

All the RAM used is dual port this greatly simplifies the design, if I had to single port RAM I would have to worry about moving data in and out of the devices and the timing of it in much greater detail, it would also slow a lot of things down. For example the CPU would either have to adopt a Harvard architecture instead of a Von Neumann or suffer a slow down.

6.7 Top level

The top level (**top_level.vhd**) is an important piece of the project, it brings all of the modules together and handles input and output (this should be moved

into a module of it's own). It also provides an interface between the outside world and the FPGA internals (ie. What the User Constraints File maps to).

Stimulus is provided for simulation purposes by **test_bench.vhd**, technically the highest up module, but only for simulation. The

6.8 Test benches and waveforms

Below is a wave form for a simple program along with said simple program further down:

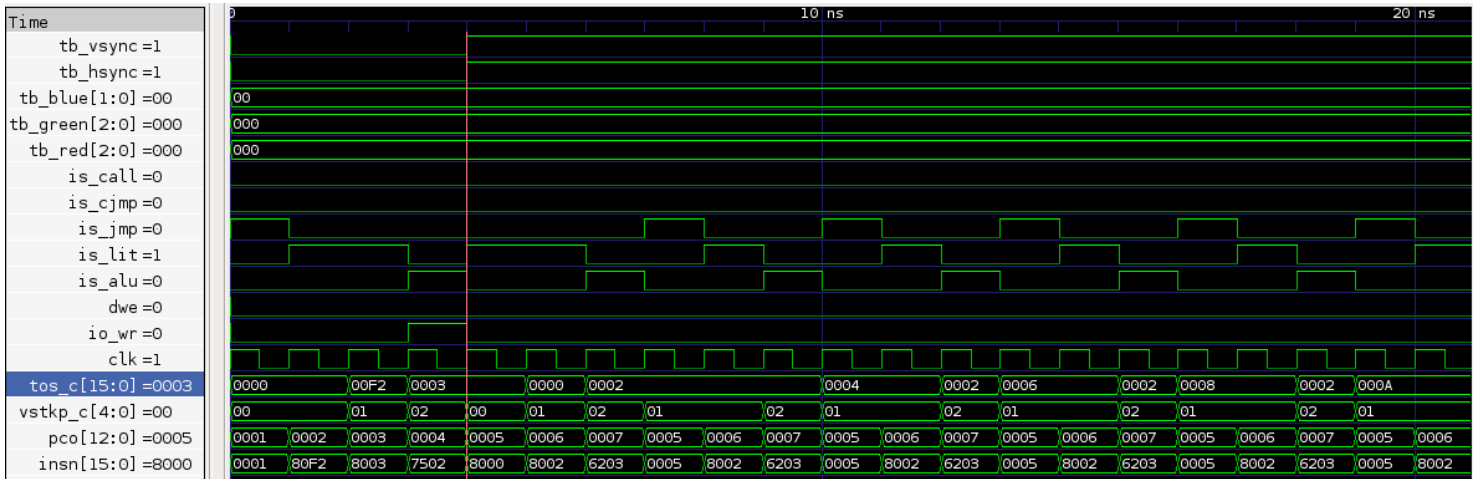


Figure 6: Adder waveform

This program sets up the VGA control register then loops around adding two to the top of the stack.

```

242 constant o_vgaCtrlDefault
3 constant o_vgaCtrl
: [SETUP]
    o_vgaCtrlDefault lit
    o_vgaCtrl lit
    _output
;
start
    [SETUP]
    0 lit
    label main
        2 lit
        _+
        main jmp
stop

```

Now given the waveform how can we show that this processor is running that code? It turns out that unless you want to do some more in depth debugging of a problem you do not really need to know what is going on in that many registers.

The following registers suffice for most needs; **tos_c**, **vstkp_c**, **pco** and **insn**. With these you can see what the current instruction being executed is each clock cycle as well as its result and effects on the variable stack pointer. (tos_c is the top of the stack, vstkp_c is the stack pointer, pco the address of the instruction to be executed and insn the instruction.

All this program does is set up the VGA control registers, push zero onto the stack and repeatedly add two to that number, we can see the VGA signals (**tb_vsync** and **tb_hsync** go high the clock cycle after the instruction "**_output**" is executed in our code (Hex number 7502). After this a zero then a two are pushed to the stack, then the two added to zero, accounting for the first long two which lasts for four clock cycles, after this is alternates between two and the current number that is to be added to. A simple test program which shows that this CPU works.

We can also see what type of instructions are being execute with **is_call**, **is_cjmp**, **is_lit** and **is_alu**, each of which are pretty self explanatory given the instruction set encoding.

Each time a new program is test, it must be first check with the test benches before going through the long process of synthesizing everything.

This one shows the same code modified to subtract instead:




Figure 7: Adder waveform

With the minor modifications:

```
start
[SETUP]
256 lit
label main
2 lit
--
main jmp
stop
```

And we can see it happily subtracting away!

6.9 Resource usage

To show you how little room the *entire* design takes up (and hence how little the CPU takes up) here is a summary a report generated by Xilinx's tools.

Device utilization summary:

Selected Device : 6slx16csg324-3

Slice Logic Utilization:

Number of Slice Registers:	303	out of	18224	1%
Number of Slice LUTs:	814	out of	9112	8%
Number used as Logic:	780	out of	9112	8%
Number used as Memory:	34	out of	2176	1%
Number used as RAM:	32			
Number used as SRL:	2			

Slice Logic Distribution:

Number of LUT Flip Flop pairs used:	879			
Number with an unused Flip Flop:	576	out of	879	65%
Number with an unused LUT:	65	out of	879	7%
Number of fully used LUT-FF pairs:	238	out of	879	27%
Number of unique control sets:	35			

IO Utilization:

Number of IOs:	46			
Number of bonded IOBs:	46	out of	232	19%

Specific Feature Utilization:

Number of Block RAM/FIFO:	12	out of	32	37%
Number using Block RAM only:	12			
Number of BUFG/BUFGCTRLs:	2	out of	16	12%
Number of DSP48A1s:	1	out of	32	3%

Of note is how few LUTs and Slice Registers are used (8% and 1% respectively), and how many BRAMs are used (12), this shows that if we wanted to we could either fit more CPU cores onto the device or substantial devices, the main limiting factor would be the BRAMs and the Slice LUTs used.

The timing, that is the longest delay path, is subject to a more frequent and sensitive changes, if I am to issue one instruction per cycle the delay needs to be minimized to less than 10^{-9} seconds. As I am currently working on the project the maximum speed is roughly 97MHz, while I was demonstrating the project it was about 102MHz which is what we want, if I want to include my new changes I will either have to accept a slow down and work around it or minimize that delay, and the delay is easy to find as the tools tell you were it is.

Bellow is a partial example of the delay estimate and route:

Timing Details:

All values displayed in nanoseconds (ns)

=====

```
Timing constraint: Default period analysis for Clock 'clk'
  Clock period: 10.293ns (frequency: 97.149MHz)
  Total number of paths / destination ports: 106202 / 825
```

```
Delay:          10.293ns (Levels of Logic = 19)
Source:         mem_h2_instance/Mram_ram8 (RAM)
Destination:    h2_instance/tos_c_15 (FF)
Source Clock:   clk rising
Destination Clock: clk rising
```

Data Path: mem_h2_instance/Mram_ram8 to h2_instance/tos_c_15

Cell:in->out	fanout	Gate Delay	Net Delay	Logical Name (Net Name)
RAMB16BWER:CLKA->DOA0	113	1.850	2.016	mem_h2_instance/Mram_ram8 (cpu_insn<14>)
LUT3:I1->0	1	0.203	0.924	h2_instance/Mmux_aluop<1>11_1 (h2_instance/Mmux_aluop<1>11)
...

This goes on longer giving you the entire route that the path takes, minimize this and you minimize the delay in your design.

One fairly trivial optimization in the design which brought resource utilization down from 23% of the slice utilization down to its current value was to properly use the distributed RAM construct, before the entire stack could be accessed all at once which was not necessary, nor possible by the current set up, only the next on stack is necessary to access, the top is held in its own separate register.

7 Problems

There were two main problems in this project, the build time for the hardware implementation and the multitude of problems faced when debugging the system, both of which slow the project down.

7.1 Build time

Some languages like C compile fairly quickly, others provide an interactive environment that you can test quickly like python. Any HDL however takes a fairly long time to simulate and a very long time to turn this design into a working 'bitfile' to be uploaded to the integrated circuit. This acts as a friction and a drag on the development process as going from a simple change in a single line of code to the hardware can take upwards of fifteen minutes (on my computer). Simulation takes up less of that time, but it is still impractical to simulate even a second of real time.

There are a few ways to mitigate this, the best one is to be right the first time when you commit to hardware, although this can be difficult, I didn't have the time to create an adequate test bench to simulate incoming UART data for example, which made debugging the UART a slow process.

The other simpler way is to use ever faster hardware, although that is an expensive option.

7.2 Debugging

Debugging can be awkward at best as there are many different places in which an error could occur and multiple errors could be conspiring to cause any given problem at hand. Why this is can be given by the list of places where a mistake could occur: The FORTH interpreter, the assembler, the assembly, the VHDL, the VHDL test-benches, the build system and finally the hardware implementation of what I have done could be different from the simulation.

The debugging process involves first pinning down where the error is as per usual, it is just made more unusually difficult by the fact that the error could be in so many different places. This is why an adequate test bench is a must.

Eventually however as the system becomes more mature (for example I worked out many of the bugs in the assembler fairly quickly) the tools used also become more stable and reliable so you do not have to worry about them as much.

8 Documentation

In this section I will provide a terse (terse in the sense that this should be its own paper) about the system.

8.1 FORTH interpreter

Instead of being a stand alone program with only one use I decided to create a full blown and reusable programming language in C, the reason for the extra complexity is because of its utility, I have used this program not only as an assembler but as part of the build process in ways that are not shown in the code (for example for converting between file formats).

The interpreter can trace its lineage back to an entry to the IOCCC ⁴. The interpreter is drastically different from the original, being a complete rewrite although there are still similarities.

The program interprets threaded code [6] for a stack machine, the C program does the initial heavy lifting allowing users to interpret commands but it is severely lacking in capabilities initially. The trick is to write most of the language *in itself*. This maybe odd to people coming from a background in a 'normal' compiled language such as C/C++ or Java. You do not modify or extend the language itself but instead provide new functionality via libraries you have written.

Initially the language does not even have basic elements such as the "if ... else ... then" statements or even loops, they are written in the language itself.

⁴ The IOCCC is the International Obfuscated C Coding Competition, the entry can be found here: <http://www.ioccc.org/all/all.tar.bz2> in the folder '1992' for 'buzzard.2'.

FORTH, and my dialect of it, has the ability that its syntax can be changed arbitrarily. Lets say you want to make an interpreter for a different language, lets say lisp, you could do that and start executing lisp, while at the same time it would be a valid FORTH program as well.

Of note is the fact that in FORTH terminology a 'Word' does not refer to a machine word but instead to a defined function, so when I refer to a 'Word' I mean function.

8.1.1 The C Program

The C program itself is design to be portable being written entirely in ANSI C [9], I have had this program running on my phone, and plan with some minor adjustments to have this running on a few embedded systems.

The C program is a threaded code interpreter and executes something which looks kind of like FORTH, although there are differences. It provides methods for extending the language, arithmetic and logic operations, conditional jumps, stack manipulations, writing to/from memory and input and output functions. It only supplies a basic set of primitives which are:

- **":"**: Read in a word and compile a header for it in the dictionary, switch to *compile* mode.
- **"immediate"**: immediate is an *immediate* word which makes the current word under going compilation *immediate*.
- **"read"**: Read in a space delimited word, if it is in the dictionary and we are in compile mode, compile a pointer in the dictionary to the found word, else execute it, else if it is a number push it onto the variable stack, else there is an error.
- **"\"**: Ignore input stream until end of line (ie. A comment).
- **"exit"**: Return from a function call.
- **"br"**: Branch unconditionally to address held in address after this instruction.
- **"?br"**: Same as "br" except branch conditionally when the top of the stack is zero.
- **"+"**: Pop two numbers off the variable stack, add them and push the result.
- **"-"**: —" —Subtract them —"
- **"*"**: —" —Multiply them —"
- **"%"**: Pop two numbers off the variable stack, perform and work out the remainder when the second item off is divided by the first, push the result.
- **"/"**: —" —, divide the second item off by the first off, push the result.
- **"lshift"**: logical shift left the next on stack by first on stack places. Push the result.

- **"rshift"**: ...same but with right shift.
- **"and"**: Pop two numbers, compute logical conjunction of them, push the result.
- **"or"**: ...same but with logical disjunction.
- **"xor"**: ...same but with exclusive disjunction.
- **"* *"**: Bitwise inversion of top of stack.
- **"1+"**: Add one to top of stack.
- **"1-"**: Subtract one from top of stack.
- **"clear"**: Clear top of stack.
- **"0>"**: Test if top of stack is less than zero, push result.
- **"="**: Pop two numbers off variable stack, test for equality, push result.
- **"<"**: Pop two items off the variable, test if the first is greater than the second, push the result.
- **">"**: ...same but the inverse.
- **"@reg"**: Use top of stack as index into the array of registers, push what is in there on to the variable stack.
- **"@dic"**: Use top of stack as index into the dictionary, push what is in there on to the variable stack.
- **"@var"**: Same but with the variable stack itself.
- **"@ret"**: Same but with the return stack.
- **"@str"**: ...with string storage.
- **"!reg"**: Store next on stack into the address pointed to by the first on the stack into the register file.
- **"!dic"**: ...into the dictionary.
- **"!var"**: ...into the variable stack.
- **"!ret"**: ...into the return stack.
- **"!str"**: ...into string storage.
- **"key"**: Push one changed from the input to the variable stack.
- **"emit"**: Pop one character from the variable stack, and output it.
- **"dup"**: Duplicate top off the variable stack.
- **"drop"**: Drop top of variable stack.
- **"swap"**: Swap top two items of variable stack.

- **"over"**: Duplicate the second item on the variable stack.
- **">r"**: Move an item from the top of the variable stack to the return stack.
- **"r>"**: ... the other way around.
- **"tail"**: Word used for recursion, call this primitive then the word you want to use recursively.
- **"'"**: Push the next compiled word onto the stack at run time.
- **","**: Advance dictionary pointer, pop the top of stack and write in into the dictionary.
- **"printnum"**: Use the top of the stack as an index into a string in string storage, treat this as a number and print it.
- **"get__word"**: Use the top of the stack as an index into string storage, store a space delimited string there.
- **"strlen"**: Pop the top of the stack and use it as an index into string storage, compute the string length and push the result.
- **"isnumber"**: Pop the top of the stack and use it as an index into string storage, test if the string is a number or not and push the result.
- **"strnequ"**: Pop two numbers off the stack and use both as indices into string storage, computer whether or not the strings are equal; zero is pushed if they are, one if they are not, two if the strings are too long ('too long' is determined in the source).
- **"__find"**: Find a word in the dictionary if it exists and push a pointer to the beginning of the word header if it does.
- **"halt"**: Halt the system.
- **"kernel"**: Use the top of the stack to perform an external call. All calls have access to all the memory.

Also defined but not present here are three 'hidden' words, that is words with no name which are:

- **"Push integer"**: Pushes the next instruction onto the data stack, advances program counter over instruction.
- **"Compile"**: Compile a pointer to the execution token 'Run' of a word.
- **"Run"**: Run a word, saving the return address onto the return stack.

The word ";" has not been defined here, it instead defined later like more of the standard words such as "rot" (rotate first three stack items).

The best documentation is the code itself, it is not too long, but I will give a quick overview of what happens.

The memory is initialized, a function that is created that calls read then calls itself (read also handles this so the stack does not blow up), this forms the basic command interpreter.

Before this is run a list of symbols is read in which forms the names of the initial dictionary. Then the function it has created is executed.

There are two states the interpreter can be in, compile or command mode, in compile mode (which is what we start out in) if a word is found a pointer to that word is compiled into the dictionary, in command mode it is executed instead. There is a certain class of words that are always executed and they are called *immediate* words.

Each word compiling has the following structure:

prev	str	compile	run	data field ...
------	-----	----------------	-----	----------------

Each immediate word has this structure:

prev	str	run	data field ...
------	-----	------------	----------------

Where "prev" is a pointer to the previous word, "str" a pointer to the words name, "compile" is the instruction compile if present and "run" is the instruction "run". The text in bold is what is run when the word is found by "read", what makes the difference between an immediate and a compiling word is what is pointed to.

The data field is of a variable length and in a normal word consists entirely or either pointers to other words execution field (either run or compile), numbers to push onto the stack or places to jump to (after either "push" or one of the branch instructions respectively).

The dictionary consists of a linked list of words, and each words data field consists of pointers to other words making for very compact code.

There are two stacks, a variable stack where computations are generally performed and a return stack for return from functions and as temporary storage. The other chunks of memory include a register file which contains pointers to the next word and instruction to execute, a pointer to the previously defined word, stack pointers, a pointer into the next available dictionary address as well as some information about the virtual machine itself. Then there is the dictionary which contains all the defined words as well as being used as a general storage facility. Finally there is a string storage space, used to store the names of words and other strings.

While this description is lacking, I do have a limited space I have to write about each of the components so for brevities sake I will have to refer you to the actual code itself.

8.1.2 Basic commands and ideas

After the basic word list has been defined we then define more FORTH words and bring the interpreter/compile into a working state, loops and conditional statements as well as words for debugging and file operations are introduced here. After this the constants and system necessary for the assembler are made, because the assembler is still executing in the FORTH environment you can still call any previously defined word there.

To interpreter is currently set up to read a start-up file, *start.fs* which then reads the assembly file *h2.fs* and outputs that to the right directory.

To run the interpreter only, you would need to edit the file *start.fs* to not run *h2.fs*, which is a trivial change at the end of the file, just delete the lines:

```
foutput ../vhdl/mem_h2.binary
finput h2.fs
```

You can then run the interpreter (./forth) and type 'words' to get a list of defined commands.

To get a better understanding of how this interpreter works consult the IOCCC submission 'buzzard.2' from 1992 [11]

A more up to date, better documented and more standalone version also written by myself is available at: <https://github.com/howej/c-forth>. This is the main branch from which this I used in this project is based.

8.2 H2 CPU

The H2 CPU has already been described, it is not too big and the VHDL in the appendix can be read and understood by someone with minimal understanding of the language.

I will instead document the assembly used.

8.2.1 Assembly

The assembler is written in a few hundred lines of FORTH, it is mostly several constants and a few helping functions specific to the H2 processor near the end of the file **start.fs**.

Bellow is a very simple test program:

```
242 constant o_vgaCtrlDefault \ 11110010, VGA control register set to this.
```

```
\ Outputs
```

```
0 constant o_7seg
1 constant o_ledS
2 constant o_vgaCursor
3 constant o_vgaCtrl
4 constant o_vgaTxtAddr
5 constant o_vgaTxtDin
6 constant o_vgaWrite
7 constant o_uartWrite
8 constant o_uartStbWrite
9 constant o_uartAckDout
```

```
\ Inputs
```

```
0 constant i_buttons
1 constant i_switches
2 constant i_vgaTxtDout
3 constant i_uartRead
4 constant i_uartAckWrite
5 constant i_uartStbDout
```

```
\ =====
\ Word definitions.
\ =====
```

```

: [SETUP]
    o_vgaCtrlDefault lit
    o_vgaCtrl lit
    _output
;

: [LED]
    o_ledS lit _output
;

: [SWITCH]
    i_switches lit _input
;

\ =====
\ Begin program loop.
\ =====
start
    [SETUP]
    label main
        [SWITCH]
        [LED]
    main jmp
stop

```

The actual program that is running on the FPGA begins at the word *start* and ends quite naturally at the word *stop*. It simply sets up the VGA to display the test pattern and after it has done that it loops around constantly read the input switches and displaying them on the output LEDs. It was the first program run on the FPGA, one which is just a simple test that everything is working.

As you can see you can call functions defined in the FORTH interpreter. In the example several words are called: '*constant*', ':' and '\'. We can define new words in the assembler to help us out with ':', simple macros in this case *[SETUP]*, *[LED]* and *[SWITCH]* that expand into the machine instructions they contain.

The program is only written out when '*stop*' is reached so messages about the program assembly progress can be written out to the standard terminal output.

Let us look at how an instruction works and how it is assembled. We can start with the simple addition instruction called *__+* in the assembler to avoid confusion with the FORTH instruction by a similar name, *+*.

The definition of this is as follows:

```

: __+ alu[ T+N d-1 or ]alu ;

```

Analysing this word by word:

- ':' : Begin the compilation of a new word in the dictionary.

- `'__+':` The name of the new word.
- `'alu[':` This is simply some syntactic sugar, it does not do anything but looks better with a corresponding `']alu'`.
- `'T+N':` This pushes the value of whatever the ALU instruction is to the variable stack.
- `'d-1':` This pushes the value of a variable stack decrement for out CPU.
- `'or':` As a stack decrement and an ALU instruction (amongst other things) can occur within the same instruction we merge these into one.
- `']alu':` This merges the instruction we want with a header (`'011'` in the high bits of the instruction) that signifies that this is an ALU operation and not for example, a conditional jump. It then writes this out into a memory model of our CPU and increments a pointer to the next available memory location in it.
- `',':` This of course ends the current word definition.

This allows us to create an assembler along with definitions along with labels which allow us to name a memory location for our program to jump to.

9 Future plans

I intend to keep working on this project after I have finished university, it offers the potential to provide a nearly complete work bench for home use and by this you will see what I mean.

There are some general improvements that could be made; making the code more uniform through out, adding variable stack sizes and word sizes, improving the instruction set, much more thorough testing (and proofs there of).

The potential improvements here are outside the scope of the project, they were never intended to be included in it as I would not have had time to do so, however when I started the project I did have these as an eventual end goal. The platform has a lot of potential.

This section fairly big despite not actually documenting the project due to the fact that it in part *justifies* the project, this section actually more or less provides the reasoning for me starting the project.

9.1 Hardware

The hardware section is the one that offers the most potential when it comes to what can be improved with the device, although it is the section that will move most slowly as HDLs can be difficult to debug, more so than normal programs.

9.1.1 Different platforms

One of the goals of this project was portable of code and not just to different Xilinx devices either. I would like to port this system to cheaper hardware such as the "Papilio One". [10] This system while less functional costs a fraction of

the Nexys 3. While the project would be more constrained, it should still be able to fit on the device with ease.

I would like to get my code running on as many platforms as possible, while it *should* be portable you can never truly find out until you have tested it on the real thing.

9.1.2 More generic code

If possible I would make my code as generic as I could. What I mean by this is being able to specify most of the parameters such as instruction size width and RAM size just by editing a few variables. Some of this might be easier than others. An example of a difficult-to-make-generic piece of code would be in the H2 CPU core:

```
1 when "01110" => tos_n <= vstkp_c & "000000" & rstkp_c;
```

Listing 1: Stack Depth Instruction

This instruction takes two stack pointers and put them on to the stack so you can analyse them. However if I wanted to increase the size of the stack the pointers would also increase in size, taking up more room and rendering this instruction incorrect.

9.1.3 PS/2 or USB keyboard

A PS/2 Keyboard could be added to the board, the Nexys 3 provides a USB to PS/2 interface, this could be used to interact with the system instead of the computer.

9.1.4 Signal Generator

A signal generator should be fairly easy to design, at least internally, the main problem would be designing the external analogue components to work from 1Hz to 100MHz (or a fraction thereof).

The VHDL would consist of a block RAM holding the data, for example a sine wave or another arbitrary signal, a counter, a pointer into the memory and a few registers to control the device. As the RAM being used is dual port, it is possible to run two signals from the same RAM at a sacrifice of precision.

This would be output in parallel over some of the boards standard I/O pins to a digital to analogue converter.

9.1.5 Data logger

By hooking up an external Analogue to Digital Converter (ADC) and a few other circuits it would be possible to have a fairly high speed data logger where plotting could be handled in real time on a normal computer (depending on how fast I manage to sample things).

This would be the most difficult to add section and it would also be the one that would provide most in the way of utility.

9.1.6 Logic Analyser

With very minimal additions a logic analyser running at the speed of the device ($\frac{1Sample}{ClockCycle}$) should be attainable, the unit does not have to be directly controlled by the CPU core, perhaps only by setting a few registers to tell it how fast to sample and how much, when to begin and end and where to send the data.

I can see the main problems being in where to store the data and how to get it off the board, I would need a faster method of communications than the UART unless I only intend to capture roughly 16 Kilobytes of data at a time.

The external hardware would be quite simple, just a way to buffer and electrically isolate the signal from the input pins of the device.

Instead of displaying the data on the device, it should be done instead on a normal desktop, at least initially.

9.1.7 Multiple cores

Due to the tiny amount of resources that are taken up by the CPU it should be possible to have more than one CPU core running on the platform at the same time, perhaps a dedicated CPU for each peripheral with one master? It would be a relatively easy way of increasing the power of the system.

This is however not a priority, first increasing functionality of the base system by adding the aforementioned peripherals would be of greater initial utility, then adding more cores would be considered.

9.1.8 Miscellaneous

The addition of timers, the reintroduction of the multiplier and adding interrupts to the new CPU core would all improve the system as a whole.

There are also several different types of memory on the Nexys 3 board which could be interfaced with which would allow me to overcome the limitations of the H2's memory restrictions, instructions could be swapped in and out.

9.2 Firmware

The firmware is what could receive the most drastic improvement, getting a small interpreter up and running on this device is a priority. This would allow it to be used as a standalone system, one which I believe would be fairly useful to multiple people.

This is where the real complexity in the device would lay, eventually.

9.3 Build system

I would like to replace most of the build system with my FORTH interpreter as I think it would be a good show case of the language I designed, it would need a way for it to deal with multiple files at the same time which it currently does not, along with a few trivial improvements, but once that has been done I could use it as a replacement for all the miscellaneous shell scripts and makefiles that I have.

10 Conclusion

In this final section I will summarize my project achievements as well as describe what I would or could have done differently if I were to do this project again.

10.1 Project achievements

From the section of "Future Plans" you can see that I have some grand ideas when it comes to what will happen with this project next and given that I would of course would like to achieve even more than I did regardless of how much I actually got done.

I did however achieve nearly all of my objectives when it came to this project and met one of the optional ones which I believe makes this an overall win.

I would have liked to have tried different instructions out with the processor and tried to optimize it for my needs, my primary concern was getting it working but there are few instructions that could be merged into one and a few others that can be replaced if for example a few of the bits were used for signed addition, bellow is an example of what I mean.

As I did not end up using the full range of the CPUs ALU bits (there are 24 instructions leaving space for further instructions, all of the bits are used, just not the full range) I could add a simple instruction that would implement signed addition. A prefix of "111" would signify that the remaining two bits should be added to the current top of the stack allowing deltas of $+1, 0, -1, -2$, as adding zero is not entirely useful this could used for something else making the most efficient use of space possible.

1	1	1	Sign	Number
---	---	---	------	--------

While this may not be the best method of doing things, I would have at least liked to have tested this out and compared the results.

I believe that my project has gone well however and has achieved its primary goals.

10.2 What I would have done differently

If I could do this project again I would have tried to make even better test benches and tried to automate even more than I did. It is possible to automatically test all the major instructions instead of manually looking at waveforms. The problem laid in the fast paced changes that I implemented, there was no use in created tests for the system if it was to change a few hours or minutes later as I was experimenting on everything.

As I have said, I would liked to have played more around with the CPU and that has already been addressed.

One more things I should have done is integrate more of other people's cores into my system, it increases functionality greatly and is not too difficult given one conditional is satisfied; *if you can find a good enough core*. Otherwise you can spend a good deal of time debugging your code and how you interface to things until you realize it is not your project that is the problem but some one else's code. More core can be found at opencores.org[12].

I would have also worked on improving the FORTH interpreter, not the C program section but the actual assembler itself, it is certain powerful enough for

the development of test programs and a bootloader but its limitation become apparent after more usage.

10.3 Final words

While giving the poster presentation a number of lecturers expressed an interest in my project and its future, specifically the capability for it to act as a signal generator or be part of some kind of DIY electronic laboratory, I am as well, so rest assured that this project will continue and check in in perhaps a year or two to monitor the progress made.

I am pleased with how this project went and I think I have managed to accomplish quite a lot; I would like to thank Marc Eberhard, my project supervisor and Kate Sugden, my course director for their help. Also I would like to thank all those people whose code has been incorporated in one form or another into this project.

11 References

References

- [1] *NexysTM3: A FPGA Spartan-6 development board*, <http://www.digilentinc.com/Products/Detail.cfm?NavPath=2,400,897&Prod=NEXYS3>, (2013)
- [2] XC6LX16 Family Overview: *Spartan-6 Family Overview*, http://www.xilinx.com/support/documentation/data_sheets/ds160.pdf, (2011)
- [3] Mark Oskin : *Advanced Digital Design: CSE 467 - Winter 2003*, <http://www.cs.washington.edu/education/courses/cse467/03wi/FPGA.pdf>, (2003)
- [4] James Bowman, Willow Garage: *J1: A small Forth CPU for FPGAs*, EuroForth 2010, <http://www.excamera.com/files/j1.pdf> (2010)
- [5] Philip J. Koopman, Jr: *Stack Computers: the new wave*, http://www.ece.cmu.edu/~koopman/stack_computers/sec6_2.html, chapter 6, section 2 (1989)
- [6] Anton Ertl: *Threaded Code*, <http://www.complang.tuwien.ac.at/forth/threaded-code.html>, (2013)
- [7] Javier Valcarce: *VHDL Macro: VGA80x40*, VHDL Monochrome VGA display adapter ,http://www.javiervalcarce.eu/wiki/VHDL_Macro:_VGA80x40 (2009)
- [8] Peter Bennett: *RS232 UART (VHDL)*, A UART core, <http://bytebash.com/2011/10/rs232-uart-vhdl/> (2012)
- [9] ANSI/ISO C: *ISO C Standards Website for the C Standard*, <http://www.open-std.org/jtc1/sc22/wg14/www/standards> (1990)
- [10] Papilio One: *A Cheap FPGA development board*, <http://papilio.cc/>,2013

- [11] IOCCC winner Buzzard: *An obfuscated FORTH interpreter*, <http://www.ioccc.org/1992/buzzard.2.design>, (1992)
- [12] Opencores: *Open source HDL cores and code*, <http://opencores.org/projects> (2013)

12 Appendix

12.1 H2.VHD

I am going to include the code for the H2 code as this forms the core of the project, for the other files go to the Github account. For instructions on how to run this consult the file "README.md" in the repository. The contents of this repository should be made available as a separate folder along with this report.

Also of note, you can look through the repository to find out how the work flow went, it does not include the beginning of the project, only about a month before the end, it does show how certain things progressed however.

```

1  -- Richard James Howe
2  -- J1 processor clone and extension. Moved bit 12 to bit 4 to
3  -- allow for more ALU instructions, added more ALU instructions.
4  -- @author          Richard James Howe.
5  -- @copyright       Copyright 2013 Richard James Howe.
6  -- @license         LGPL
7  -- @email           howe.r.j.89@gmail.com
8  library ieee, work, std;
9  use ieee.std_logic_1164.all;
10 use ieee.numeric_std.all;
11
12 entity h2 is
13     port(
14         clk:         in  std_logic;
15         rst:         in  std_logic;
16         -- IO interface
17         io_wr:       out std_logic;
18         io_din:      in  std_logic_vector(15 downto 0);
19         io_dout:     out std_logic_vector(15 downto 0);
20         io_daddr:    out std_logic_vector(15 downto 0);
21         -- RAM interface
22         pco:         out std_logic_vector(12 downto 0);
23         insn:        in  std_logic_vector(15 downto 0);
24
25         dwe:         out std_logic; -- data write enable, read
26                     enable not need.
27         din:         in  std_logic_vector(15 downto 0);
28         dout:        out std_logic_vector(15 downto 0);
29         daddr:       out std_logic_vector(12 downto 0)
30     );
31 end;
32
33 architecture rtl of h2 is
34
35     -- Program counter.
36     signal pc_c      : std_logic_vector(12 downto 0) := (others
37         => '0');
38     signal pc_n      : std_logic_vector(12 downto 0) := (others
39         => '0');

```

```

type stk is array (31 downto 0) of std_logic_vector(15
    downto 0);
39 -- Variable stack (RAM Template)
signal vstkp_c : std_logic_vector(4 downto 0) := (others =>
    '0');
41 signal vstkp_n : std_logic_vector(4 downto 0) := (others =>
    '0');
signal vstk_ram : stk := (others => (others => '0'));
43 -- Return stack (RAM Template)
signal rstkp_c : std_logic_vector(4 downto 0) := (others =>
    '0');
45 signal rstkp_n : std_logic_vector(4 downto 0) := (others =>
    '0');
signal rstk_ram : stk := (others => (others => '0'));
47 -- Stack deltas
signal dd : std_logic_vector(4 downto 0) := (others =>
    '0');
49 signal rd : std_logic_vector(4 downto 0) := (others =>
    '0');

51 -- is_x signals, booleans, does the instruction have a certain
    property.
signal is_alu : std_logic := '0';
53 signal is_lit : std_logic := '0';
signal is_jump : std_logic := '0';
55 signal is_cjmp : std_logic := '0';
signal is_call : std_logic := '0';

57 -- Top of stack, and next on stack.
59 signal tos_c : std_logic_vector(15 downto 0) := (others =>
    '0');
signal tos_n : std_logic_vector(15 downto 0) := (others =>
    '0');
61 signal nos : std_logic_vector(15 downto 0) := (others =>
    '0');
-- Top of return stack.
63 signal rtos_c : std_logic_vector(15 downto 0) := (others =>
    '0');

65 -- aluop is what is fed into the alu.
signal aluop: std_logic_vector(4 downto 0) := (others =>
    '0');
67 -- pc_plus_1, forces fewer adders.
signal pc_plus_one: std_logic_vector(12 downto 0) := (others =>
    '0');
69 -- Stack signals
signal dstkW: std_logic := '0';
71 signal rstkD: std_logic_vector(15 downto 0) := (others =>
    '0');
signal rstkW: std_logic := '0';
73 begin
    -- is_x
75 is_alu <= '1' when insn(15 downto 13) = "011" else '0';
is_lit <= '1' when insn(15) = '1' else '0';
77 is_jump <= '1' when insn(15 downto 13) = "000" else '0';
is_cjmp <= '1' when insn(15 downto 13) = "001" else '0';
79 is_call <= '1' when insn(15 downto 13) = "010" else '0';

81 -- Stack assignments
nos <= vstk_ram(to_integer(unsigned(vstkp_c)));
83 rtos_c <= rstk_ram(to_integer(unsigned(rstkp_c)));

```

```

85  -- I/O assignments
      pco      <= pc_n;
87  dout      <= nos;
      daddr    <= tos_c(12 downto 0);
89  dwe      <= insn(5) when is_alu = '1' else '0';

91  -- io_wr are handled in the ALU,
92  -- this makes things slower but we have
93  -- run out of instruction bits to use.
      io_dout   <= nos;
95  io_daddr  <= tos_c;

97  -- misc
      pc_plus_one <= std_logic_vector(unsigned(pc_c) + 1);

99  -- Signed addition!
101 dd <= insn(1) & insn(1) & insn(1) & insn(1) & insn(0);
      rd <= insn(3) & insn(3) & insn(3) & insn(3) & insn(2);

103 dstkW <= '1' when is_lit = '1' or (is_alu = '1' and insn(7) =
      '1') else '0';

105
      stackWrite: process(
107         clk
      )
109 begin
      if rising_edge(clk) then
111         if dstkW = '1' then
            vstk_ram(to_integer(unsigned(vstkp_n))) <=
                tos_c;

113         end if;

115         if rstkW = '1' then
            rstk_ram(to_integer(unsigned(rstkp_n))) <=
                rstkD;

117         end if;
      end if;
119 end process;

121 alu_sel: process(
      insn
123 )
      begin
125         case insn(14 downto 13) is
            when "00" => aluop <= "00000"; -- ubranch
127             when "01" => aluop <= "00000"; -- call
            when "10" => aluop <= "00001"; -- 0branch
129             when "11" => aluop <= insn(12 downto 8); -- alu
                operation.
            when others => aluop <= "XXXXX";

131         end case;
      end process;

133
      -- ALU
135 alu: process(
      is_lit ,
137     tos_c, nos, rtos_c,
      din, insn, aluop,
139     io_din,
      vstkp_c, rstkp_c
141 )
      begin

```

```

143 io_wr      <= '0';
144   tos_n    <= tos_c;
145   if is_lit = '1' then
146       tos_n <= "0" & insn(14 downto 0);
147   else
148       case aluop is -- ALU operation, 12 downto 8
149           -- Original J1 instructions --
150           when "00000" => tos_n <= tos_c;
151           when "00001" => tos_n <= nos;
152           when "00010" =>
153               tos_n <= std_logic_vector(unsigned(tos_c)+
154                   unsigned(nos));
155           when "00011" => tos_n <= tos_c and nos;
156           when "00100" => tos_n <= tos_c or nos;
157           when "00101" => tos_n <= tos_c xor nos;
158           when "00110" => tos_n <= not tos_c;
159           when "00111" =>
160               if
161                   nos = tos_c
162               then
163                   tos_n <= (0 => '1', others
164                       => '0');
165               else
166                   tos_n <= (others => '0');
167               end if;
168           when "01000" =>
169               if
170                   signed(nos) < signed(tos_c)
171               then
172                   tos_n <= (0=>'1', others =>
173                       '0');
174               else
175                   tos_n <= (others => '0');
176               end if;
177           when "01001" =>
178               tos_n <= std_logic_vector(unsigned(nos) srl
179                   to_integer(unsigned(tos_c(3 downto 0))));
180           when "01010" => tos_n <= std_logic_vector(
181               unsigned(tos_c)-1);
182           when "01011" => tos_n <= rtos_c;
183           when "01100" => tos_n <= din;
184           when "01101" =>
185               tos_n <= std_logic_vector(unsigned(nos) sll
186                   to_integer(unsigned(tos_c(3 downto 0))));
187           when "01110" => tos_n <= vstkp_c & "000000" &
188               rstkp_c; -- depth of stacks
189           when "01111" =>
190               if
191                   nos < tos_c
192               then
193                   tos_n <= (0=>'1', others =>
194                       '0');
195               else
196                   tos_n <= (others => '0');
197               end if;
198           -- Additional instructions --
199           when "10000" => tos_n <= std_logic_vector(
200               unsigned(nos)-unsigned(tos_c));
201           when "10001" => tos_n <= tos_c xnor nos;
202           when "10010" => -- Reserved for multiplier low
203               bits
204           when "10011" => -- Reserved for multilpier high

```



```

195         bits
        when "10100" => tos_n <= io_din; -- Should be
            integrated din instruction
        when "10101" => io_wr <= '1'; -- Should be
            integrated to other write instruction
197        when "10110" =>
            tos_n <= std_logic_vector(unsigned(nos) ror
                to_integer(unsigned(tos_c(3 downto 0))));
199        when "10111" =>
            tos_n <= std_logic_vector(unsigned(nos) rol
                to_integer(unsigned(tos_c(3 downto 0))));
201        when "11000" => tos_n <= (others => '0');
        when "11001" =>
203        when "11010" =>
        when "11011" =>
205        when "11100" =>
        when "11101" =>
207        when "11110" =>
        when "11111" =>
209        when others => tos_n <= (others => 'X');
        end case;
211    end if;
end process;

-- Reset and state-machine clock.
215 nextState: process(clk, rst)
begin
217    if rst='1' then
        vstkp_c <= (others => '0');
219        rstkp_c <= (others => '0');
        pc_c <= (others => '0');
221        tos_c <= (others => '0');
        elsif rising_edge(clk) then
223            vstkp_c <= vstkp_n;
            rstkp_c <= rstkp_n;
225            pc_c <= pc_n;
            tos_c <= tos_n;
227        end if;
end process;

229
mainProcess: process(
231    pc_c,
    insn,
233    vstkp_c, vstk_ram, dd,
    rstkp_c, rstk_ram, rd,
235    tos_c,
    is_jump, is_cjmp, is_call, is_lit, is_alu,
    pc_plus_one
237)
begin
239    vstkp_n <= vstkp_c;
    rstkp_n <= rstkp_c;
    -- main control
243    if is_lit = '1' then
        vstkp_n <= std_logic_vector(unsigned(vstkp_c)+1);
245        rstkW <= '0';
        rstkD <= "000" & pc_plus_one;
247    elsif is_alu = '1' then
        rstkW <= insn(6);
249        rstkD <= tos_c;
        -- Signed addition.
251        vstkp_n <= std_logic_vector(unsigned(vstkp_c) +

```

```

253         unsigned(dd));
        rstkp_n <= std_logic_vector(unsigned(rstkp_c) +
        unsigned(rd));
254     else
255         if is_cjmp = '1' then
            vstkp_n <= std_logic_vector(unsigned(vstkp_c) - 1)
            ;
        end if;
256
257         if is_call = '1' then
259             rstkp_n <= std_logic_vector(unsigned(rstkp_c) + 1)
            ;
            rstkW    <= '1';
            rstkD    <= "000" & pc_plus_one;
        else
263             rstkW    <= '0';
            rstkD    <= "000" & pc_plus_one;
265         end if;
        end if;
267     end process;

269     pcUpdate: process(
        pc_c, insn, rtos_c, pc_plus_one, tos_c,
271         is_jump, is_cjmp, is_call, is_alu
    )
273     begin
        pc_n    <= pc_c;
275         if is_jump = '1' or (is_cjmp = '1' and tos_c = X"0000") or
            is_call = '1' then
            pc_n    <= insn(12 downto 0);
277         elsif is_alu = '1' and insn(4) = '1' then
            pc_n    <= rtos_c(12 downto 0);
279         else
            pc_n    <= pc_plus_one;
281         end if;
        end process;
283     end architecture;

```

Listing 2: Stack Depth Instruction

12.2 Original specification

Here is the original specification for the project which can be found in the repository in full also under the name "spec.pdf".

12.2.1 FPGA: Main plan.

As I have said, this would be my main plan, to use an FPGA to make my own computing system, this is not infeasible and has done before, my own design will be unique and customized to the task I want it to do.

While it will not be useful as a general purpose computer, it could be used for pedagogical purposes, both for me and for other people as well.

The hardware would be very simple, and at minimum would be a Nexys 3 development board, with a serial connector interface. The FGPA contained in the device is a Xilinx Spartan-6 XC6SLX16 which is enough to fit my project on along with enough internal RAM to run its program.

Deliverables

The following are a list of the projects deliverables.

- CPU and Waveforms.
- Cross assembler.
- Serial Interface.
- FORTH language interpreter.
- (Optional): Keyboard Interface.
- (Optional): VGA Interface.

The last two being entirely optional, it all depends on how feasible it is to make the last two.

Languages to use

The following is the languages I wish to use, for a variety of tasks, from hardware description, to automation.

- VHDL: Hardware description of CPU core and peripherals.
- VHDL: Test Benches for hardware.
- Python: Cross Assembler.
- Assembly: Used to write the FORTH interpreter in.
- FORTH: Ultimate goal of the project, a FORTH language interpreter.

Other miscellaneous languages used will be \LaTeX for writing up the report, GNU Make files for the build system, and Bash shell scripts for automating simple tasks. The Git version control system will be used to keep track of the latest build and documentation.

Industry and other standards

As I have said I will be linking to other standards, which I will use in my project, some of which will be part of optional modules such as the Keyboard and VGA (video) output.

- J1 Processor Core: Described in a paper 4 and written in Verilog, this will be rewritten in VHDL and modified.
- Serial Port: This is a very well known standard, and will form the primary interface between the outside world and my CPU core.
- VGA standard : A well known Video standard (optional)
- PS/2 Keyboard (Optional)

Software to use

I already have the software and build environment set up and ready to be used, I will have to adapt some Makefiles for my purpose, but apart from this, the software environment I will be working on is ready. The project will be developed in a Linux environment (Debian 6.0), using Xilinx Webpack (Latest version), and Digilents drivers. Other software that is relevant to the project is Git, for versioning, GNU make, and finally, GHDL for simulation purposes only.

J1 Processor

Instead of specifying my own instruction set, I am going to use one that is already specified for me, as my device does not have to fit into such a small space, I can utilize much more of the chip if I want to, the original processor had to fit inside a camera's FPGA that had many other tasks to do.

I could experiment with optimizing the core for my specific needs, for example, I could decide if I want to use unsigned or signed arithmetic, I could also add more operations to the core (one bit is unused in the instruction set). I am refraining to say as to what I may add, I would like to see what would be the bottlenecks in my design before I start to customize the core. A big boost would probably be having a general purpose multiply instruction, however I may be able to get away with a left shift for multiplication (by 2^n only), and division by 2^n as well (right shift).

Depending on what I want to do, I could add instructions for cyclic shifts, or even more complicated things instructions tailored to my task. This is optional.

Error handling and multitasking (Optional).

I would like to create a robust system, and I would like to do that on many levels, testing as much as possible and eliminating sources of error, in that vein, there could be certain systems in place for checking whether certain things make sense. Bounds checking could be enforced at the hardware level if needed, making sure certain instructions can not be executing if they operate on data in off limits areas.

Another possibility is to aid multitasking, either with multiple cores, or with banks of registers that can be switched to and from in a few instruction cycles. Both of these tasks would be entirely optional.

Again, the J1 Processor can be found here <http://www.excamera.com/sphinx/fpga-j1.html> and here <http://www.excamera.com/files/j1.pdf>.